

Cloud Foundry Container Runtime

Lab 10 - Kubernetes Exploration

Kubernetes clusters track and manage objects of various "kinds". Applications make use of four kinds of objects in particular:

- **Pods** – groups of containers deployed as a unit
- **Replica Sets** – sets of pods defined by a template which the Controller Manager replicates across the cluster
- **Deployments** – a rollout strategy for pods and replica sets
- **Services** – end points used to distribute requests to one of a pod's replicas

Thus basic Kubernetes applications consist of pods, which implement the application functionality; replica sets, which ensure pods are always available; and Services which expose a dynamic set of pods to clients as a single endpoint. deployments describe how to launch or upgrade a given application.

1. kubectl

The `kubectl` command provides a range of features we can use with Kubernetes. Run `kubectl` without arguments to get a list of the available commands.

```
user@ubuntu:~$ kubectl
```

```
kubectl controls the Kubernetes cluster manager.
```

```
Find more information at: https://kubernetes.io/docs/reference/kubectl/overview/
```

Basic Commands (Beginner):

<code>create</code>	Create a resource from a file or from stdin.
<code>expose</code>	Take a replication controller, service, deployment or pod and expose it as a new Kubernetes Service
<code>run</code>	Run a particular image on the cluster
<code>set</code>	Set specific features on objects
<code>run-container</code>	Run a particular image on the cluster. This command is deprecated, use "run" instead

Basic Commands (Intermediate):

<code>get</code>	Display one or many resources
<code>explain</code>	Documentation of resources
<code>edit</code>	Edit a resource on the server
<code>delete</code>	Delete resources by filenames, stdin, resources and names, or by resources and label selector

Deploy Commands:

<code>rollout</code>	Manage the rollout of a resource
<code>rolling-update</code>	Perform a rolling update of the given ReplicationController
<code>scale</code>	Set a new size for a Deployment, ReplicaSet, Replication Controller, or Job
<code>autoscale</code>	Auto-scale a Deployment, ReplicaSet, or ReplicationController

Cluster Management Commands:

<code>certificate</code>	Modify certificate resources.
<code>cluster-info</code>	Display cluster info
<code>top</code>	Display Resource (CPU/Memory/Storage) usage.
<code>cordon</code>	Mark node as unschedulable
<code>uncordon</code>	Mark node as schedulable
<code>drain</code>	Drain node in preparation for maintenance
<code>taint</code>	Update the taints on one or more nodes

Troubleshooting and Debugging Commands:

<code>describe</code>	Show details of a specific resource or group of resources
<code>logs</code>	Print the logs for a container in a pod
<code>attach</code>	Attach to a running container
<code>exec</code>	Execute a command in a container
<code>port-forward</code>	Forward one or more local ports to a pod
<code>proxy</code>	Run a proxy to the Kubernetes API server
<code>cp</code>	Copy files and directories to and from containers.
<code>auth</code>	Inspect authorization

Advanced Commands:

<code>apply</code>	Apply a configuration to a resource by filename or stdin
<code>patch</code>	Update field(s) of a resource using strategic merge patch
<code>replace</code>	Replace a resource by filename or stdin
<code>convert</code>	Convert config files between different API versions

Settings Commands:

<code>label</code>	Update the labels on a resource
<code>annotate</code>	Update the annotations on a resource

```
completion      Output shell completion code for the specified shell (bash or zsh)

Other Commands:
  api-versions   Print the supported API versions on the server, in the form of "group/version"
  config         Modify kubeconfig files
  help           Help about any command
  plugin         Runs a command-line plugin
  version        Print the client and server version information

Usage:
  kubectl [flags] [options]

Use "kubectl <command> --help" for more information about a given command.
Use "kubectl options" for a list of global command-line options (applies to all commands).
user@ubuntu:~$
```

Take a moment to review available options. One useful subcommand is the global options, take a moment to review the output of `kubectl options`.

To use the `kubectl` command to control a remote cluster we must specify the cluster endpoint to `kubectl`. The `kubectl` command can be used to control several clusters from a single workstation. Clusters are given a name and settings, including the IP address and port of the cluster API service.

To get configuration help issue the `kubectl help` subcommand.

```
user@ubuntu:~$ kubectl help config
```

Modify kubeconfig files using subcommands like "kubectl config set current-context my-context"

The loading order follows these rules:

1. If the `--kubeconfig` flag is set, then only that file is loaded. The flag may only be set once and no merging takes place.
2. If `$KUBECONFIG` environment variable is set, then it is used a list of paths (normal path delimitting rules for your system). These paths are merged. When a value is modified, it is modified in the file that defines the stanza. When a value is created, it is created in the first file that exists. If no files in the chain exist, then it creates the last file in the list.
3. Otherwise, `$(HOME)/.kube/config` is used and no merging takes place.

```
Available Commands:
  current-context  Displays the current-context
  delete-cluster   Delete the specified cluster from the kubeconfig
  delete-context   Delete the specified context from the kubeconfig
  get-clusters     Display clusters defined in the kubeconfig
  get-contexts     Describe one or many contexts
  rename-context   Renames a context from the kubeconfig file.
  set              Sets an individual value in a kubeconfig file
  set-cluster      Sets a cluster entry in kubeconfig
  set-context      Sets a context entry in kubeconfig
  set-credentials  Sets a user entry in kubeconfig
  unset            Unsets an individual value in a kubeconfig file
  use-context      Sets the current-context in a kubeconfig file
  view             Display merged kubeconfig settings or a specified kubeconfig file
```

```
Usage:
  kubectl config SUBCOMMAND [options]
```

Use "kubectl <command> --help" for more information about a given command.
Use "kubectl options" for a list of global command-line options (applies to all commands).
user@ubuntu:~\$

Run the `kubectl config view` subcommand again to display the current client configuration.

```
user@ubuntu:~$ kubectl config view

apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: REDACTED
    server: https://172.16.151.229:6443
    name: kubernetes
```

```

contexts:
- context:
  cluster: kubernetes
  user: kubernetes-admin
  name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED
user@ubuntu:~$

```

When you run *kubectl* commands a context is required. The context tells *kubectl* which cluster to connect to and which user to authenticate as. As you can see the values kubeadm configured means the *kubectl* command tries to reach the API server on port 6443 via our host's IP with TLS.

To view the REDACTED elements, add `--flatten`.

We can configure *kubectl* explicitly so that we can adjust our cluster settings in the future if need be. Get help on the `config set-cluster` subcommand:

```

user@ubuntu:~$ kubectl help config set-cluster

Sets a cluster entry in kubeconfig.

Specifying a name that already exists will merge new fields on top of existing values for those fields.

Examples:
# Set only the server field on the e2e cluster entry without touching other values.
kubectl config set-cluster e2e --server=https://1.2.3.4

# Embed certificate authority data for the e2e cluster entry
kubectl config set-cluster e2e --certificate-authority=~/.kube/e2e/kubernetes.ca.crt

# Disable cert checking for the dev cluster entry
kubectl config set-cluster e2e --insecure-skip-tls-verify=true

Options:
  --certificate-authority='': Path to certificate-authority file for the cluster entry in kubeconfig
  --embed-certs=false: embed-certs for the cluster entry in kubeconfig
  --insecure-skip-tls-verify=false: insecure-skip-tls-verify for the cluster entry in kubeconfig
  --server='': server for the cluster entry in kubeconfig

Usage:
  kubectl config set-cluster NAME [--server=server] [--certificate-authority=path/to/certificate/authority]
[--insecure-skip-tls-verify=true] [options]

Use "kubectl options" for a list of global command-line options (applies to all commands).
user@ubuntu:~$

```

kubectl configuration data is saved in a YAML file in your `$HOME/.kube` directory using these commands. Display the configuration file we copied in lab 2.

```

user@ubuntu:~$ ls -la ~/.kube/

total 24
drwxrwxr-x  4 user user 4096 Dec  4 11:30 .
drwxr-xr-x 17 user user 4096 Dec  4 12:57 ..
drwxr-xr-x  3 user user 4096 Dec  4 11:30 cache
-rw-----  1 user root 5454 Dec  4 11:29 config
drwxrwxr-x  3 user user 4096 Dec  4 12:20 http-cache
user@ubuntu:~$

```

Display the contents of the config file:

```

user@ubuntu:~$ cat ~/.kube/config

apiVersion: v1
clusters:
- cluster:
  certificate-authority-data:

```

```
server: https://172.16.151.229:6443
name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tCk1JSUM4akNDQWRxZ0F3SUJBZ0l1JVjr1Q056a2RSUEV3RFFZSktvWkloMNOQVFFTEJRCXQdGVEVUTUJF
R0ExVUUkQXhNS2EzVmIawEp1wIhSbGn6QWVGdzB4TnpFeU1EUxhpVEV4TURKYUZ3MHHhPREV5TURReE9URXhNRE5hTURReApGekFWQmDOvQJBb1REBk41
YzNSbGJGUChRZWE4wW1hKek1Sa3dGd1lEVlFRREV4QnJkV0psY201bGRHVnpMV0ZrCmJXbHVNSU1CSWpBTkNa3Foa2lHOXcwQkFRRUZBQU9DQVE4QUU1J
SUJZD0tDQVFFQXdzZ0FUQzdLa0szU1hMRWEKewp1eG43S2pVR21LV3RwN1VTTnJnWEk5Ww1kNXdwZWRNMct2amttSGM3WjgzYTJrUFG1SExG0EdwcEFi
S2FNTgpQS1Vcy205ekRKZk10Q1FKWFBhavlUkFET3hRTMTPVPL0ZiWgDLTG85ZDh1ay9McM5jZFUybnBhc1h3dnFmdnlxCitvVdFQ3kvaTJHL0V0NkRk
NlJ0cmIybgGk4MD1YSXpQZW4xczQxUkZ1Yko5T01vendBa29FdkJRORUR3RC9EUUwKawXVeGwXOFY5WVhMakVFQmFES2MxR3NSOThqU1ZJRmVKT29TaWUx
QmYveTJaMHphVHRHbzhKc3Mrb1JKdmVrbAorMGZLYU1UcGxRYURQWmR6OFcrZWwVUZsd1h3RTFTbi9JcDVxNVJMdlh3TVB0Qkx4Tk1GdXc5bjZzcXM0
Mm9vCnB5MlNjd0lEQVFBQm95Y3dkVEFPQmDOvQkhROEJBZjhFQkFNQ0JhQXdFdl1EVl1IwEJBd3dDZ1lJS3dZQkJRUVuGkQXJd0RRWUlpLb1pJaHJzTkFR
RUXCUUFEZ2dFQkFJQk1qV3lnbUtlVkl0Mk5MRkxvQnh5UDRYVExwRmZWNPaeQp2Y1c1SmcxQXNUUGZFQmcxSEhKRk8xb2g2L21rQ05PTnZESXN3bDR2
bi9UWw1ZlZlJBQ09pUjNHeDkrYnZLcU14CkYydl1taItbnbXFcvc0Nqa0ZDZGFudjE3RzNwVFJicjhrQndrVnpKeStIL1dzWUtpMkIxcHdId3hCazJkcm1X
ZlIKN0ErTFd5VU5uamNnNwdiam83d1FObwIzTG9zMVJMMys4UU1scFlCTjMrdfZVL1pON3l3NG9LVGTFeDNOR1RlSAPXaE9kT0xaTzAxSzhIwWtqQ2hh
SXRTDl1Ja0I4ew5M50lRcXU2UndYclJ0bkRrd1F1MDZTY0Q4dW4xS0VkdUs0CkVrMnFmdW0vdHprN0U2aEFBcEhtM1dalZBxc5m5JU0h0Y2U1cE93Q2Rq
MHRJT255Y1NRUT0KLS0tLS1FtQgQ0VSVElGSUNBVEUtLS0tLQo=
```

L\$0tLS1CRUDJTtBSU0EgUFJJVKFURSBLRVktLS0tLQpNSUIFcEFJQkFBS0NBuUVbd3NnQVRDn0trSzNSWExFYXlqdXhuN0tqVudtS1d0cDZVU05yZ1hJ
OV1pZDV3VmVkJCk0wK3Zq21IYzdaODNhmMtQWDVITEY4R3BwQWJLYU10UEpVQmNtOXpESmZNdENRSlhQVWlZbKJBRE94UTE1Ty8KRKhYz0tMbZ1kOHVr
L0xybmNkVTJucGFyWHD2cWZ2EXerb1VXRUN5L2kyRy9FdBZEZDZSVHJiMmxp0DA5WE16UAp1bJfZNDfSRnV1Sj1PTW96d0Frb0V2Q1E5RHdEL0RRTG1s
VXhSMThWOV1YTGpFRUJhREtjMUdzUjK4a1JWSUZlCkpPb1NpZTFZCi95MlowemFudEdv0EpzcytvUkpZDZwtsKzBmS2FNVBHSUWFEK1pkejHXK2VkcFVG
bHdYd0UxU24KL0lWNXE1Ukx2WHDuNHURCHthOTUZ1dZluNnNxczQyb29wE1TjY3dJREFRQUJBB0lCQVFDTWfudjBiNkx0NjdCTApQdzJPRVh4ODROM2s0
VUN4UuUE1R5WjRHZS93YTM3VmQ2gZtXzEkE4SS9M1p1nPTM3RZd1Q2t1W0N9YVSgXlMnZGcmQXcmfLcHt0YjRdajZjTlPbW3ZCcmJiZ2JNwVmlpNczNhHpY
VHZTCe1rZrJ20X1uSVc4RHZpZ0luRDMWwk50RGsKbkktkeGxER1VsWDI4TUVuN2cxZUpEM2lYZm1qeVAwNEXFR1BkbFNyRVmXbExlNHN1R0RJUT1VTHVv
anpZrkFqYgpnEudEUzIwVTZ5N2Fec0lZdVdEQ2ZhdEx2Ym5mQy9VZSsvYnN6b1Q3dEdtRzLzY1XL013S21UQmI0MmlaeFhECnN2RUNXY1J2c2FpAVBm
ckJ2bk41a3M4enhRQk1lWlVYRnlzRmUxc09bk3Q0Tms5Z29MU2RZRk1CaDBRV3drKy8KZnRUSGZUa0JBB0dCQU1oWGI0cG8xZGhubUw1VnpXWlZiBUFS
cEU3STRGR1ZHY0F1VXNCczUvaFI5UGhhNE5yKwp6MlFWSkTBK05ZYlJYZThPdWxtdzRZeUpaUFN2RVBFYjlmVG1pY091c1ZoMHBCSHlzYnZEY0R5TjdH
c1NLajMzCksyQzV0BgHFTFo2YmJETz1LamU0cFNWazZWUtDVeHlQUHJyRE4rdGhvRmNDdWRsTzBlcFhRQUtCQW9HQkFQamwKSENZR3NyMk1Md3RFb1pW
Qy9MS3orU2t1WU0wekRkZkhtVNMXXY1VUem90Bg5vUwHLN2FvYzZscjRjVEYvTkZBSQPbC0NMS3k3e1JzY05LY1M1UFdwQ1VMTjJlMSvHhZEsRGFT1Fb
N0NCSUJBK1V4SHU0eHA4UEFNGKX13dUthHCmoyM1p0aXlJ5jJQjUd1VYvUwNDBWTWIScjJEdEk4b3pid09pQjdtYlPbW3ZCcmJiZ2JNwVmlpNczNhHpY
Y3FtOTUKWDVRd0lpWU13Yi9uQ3dnUCtHtm5XekhOekY5a2tONTZFAnNRbVk0ZDJZNDQ2Vy8vcmZnemtcjv1rMUZZbWl0agpLT3IwwjczZjFJTHpYeXlJ
K2E1QVl1v01zaD14RGk0aDlJQXdkRF1vZ25pLzG1NGcyM2pFK2xcVGtKaDBQWkv5CkFxeFRNWHB4ZUZ3R0VYOTRzM3dDT1FBQkFvR0FNS0ZyVjQ2MWU4
e1JERTJUbUx0btDtktZf2ak11bk5sZDJneFkKekt0STuYvUHLNtFRSU9EckfQTDNZMkRwbFBWR2phQ1VVU1NnRW10Y0wrWkZnTmMweGI5Q1QyQ2t3MXfj
VC9PUQp1UFdaa3ZJWDFyU212SGxGakZaQ0gyaDlkaJNaMlHlNXNZZjVUVWdwRWWhyaHA0YnJ5NkFaz1VKTUZJRlRtDxh0CkM4emZRcnNDZ1lCTExvtJrI
M0xakVORyt5K2tFtFmcEp4MDJhNk4zeUN5SmFyNwIzwbVjS0YvMG5nTm9PZlCkCufKRWE3ZzFdaVvKMVJibGtZZHxhdFnIYUZEuZVERDI5MHVLOEzj
WEVRDWRlCEJ1Yk1LDQ0oZc3ZWdndIRElHegphdm8zc3lSd3lVdGt0MmZmYk0rbHJETE8zTmg3RnJiS3NOWG5zTjgrYjvBeDBZSHhkMXRsUUE9PQotLS0t
LUVORC0R0U0EgUFJJVKFURSBLRVktLS0tLQo=

user@ubuntu:~\$

2. Test the Cluster

```
user@ubuntu:~$ kubectl cluster-info
```

Page 4/15

```
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
user@ubuntu:~$
```

If you are really adventurous run the suggested command for a detailed cluster overview, careful though, its a lot of information!

```
user@ubuntu:~$ kubectl cluster-info dump |& wc -l

5215
user@ubuntu:~$
```

To get detailed node information use the *describe node* subcommand again on the desired node name:

```
user@ubuntu:~$ kubectl describe node ubuntu

Name:          ubuntu
Roles:         master
Labels:        beta.kubernetes.io/arch=amd64
               beta.kubernetes.io/os=linux
               kubernetes.io/hostname=ubuntu
               node-role.kubernetes.io/master=
Annotations:   node.alpha.kubernetes.io/ttl=0
               volumes.kubernetes.io/controller-managed-attach-detach=true
CreationTimestamp: Wed, 28 Mar 2018 10:14:15 -0700
Taints:        <none>
Unschedulable: false
Conditions:
  Type             Status  LastHeartbeatTime             LastTransitionTime             Reason
  ----             -
  OutOfDisk        False   Wed, 28 Mar 2018 11:44:40 -0700   Wed, 28 Mar 2018 10:14:12 -0700
  KubeletHasSufficientDisk   kubelet has sufficient disk space available
  MemoryPressure   False   Wed, 28 Mar 2018 11:44:40 -0700   Wed, 28 Mar 2018 10:14:12 -0700
  KubeletHasSufficientMemory kubelet has sufficient memory available
  DiskPressure     False   Wed, 28 Mar 2018 11:44:40 -0700   Wed, 28 Mar 2018 10:14:12 -0700
  KubeletHasNoDiskPressure  kubelet has no disk pressure
  PIDPressure      False   Wed, 28 Mar 2018 11:44:40 -0700   Wed, 28 Mar 2018 10:14:12 -0700
  KubeletHasSufficientPID   kubelet has sufficient PID available
  Ready            True    Wed, 28 Mar 2018 11:44:40 -0700   Wed, 28 Mar 2018 11:07:18 -0700   KubeletReady
kubelet is posting ready status. AppArmor enabled
Addresses:
  InternalIP: 192.168.225.210
  Hostname:   ubuntu
Capacity:
  cpu:                2
  ephemeral-storage:  18447100Ki
  hugepages-1Gi:      0
  hugepages-2Mi:      0
  memory:             2029876Ki
  pods:              110
Allocatable:
  cpu:                2
  ephemeral-storage:  17000847332
  hugepages-1Gi:      0
  hugepages-2Mi:      0
  memory:             1927476Ki
  pods:              110
System Info:
  Machine ID:         6e883acc04fc7db3713776be57a3dac9
  System UUID:        B2344D56-FE13-D736-51F0-31DF1E2E8A04
  Boot ID:            9eda422d-14ef-43aa-94d7-5da65fbc2dfa
  Kernel Version:     4.4.0-116-generic
  OS Image:           Ubuntu 16.04.1 LTS
  Operating System:   linux
  Architecture:       amd64
  Container Runtime Version: docker://18.3.0
  Kubelet Version:     v1.10.0
  Kube-Proxy Version:  v1.10.0
ExternalID:          ubuntu
Non-terminated Pods: (7 in total)
  Namespace           Name
  ----             -
  Limits
  -----
  ---

```

```

kube-system          etcd-ubuntu          0 (0%)    0 (0%)    0 (0%)    0 (0%)
kube-system          kube-apiserver-ubuntu 250m (12%) 0 (0%)    0 (0%)    0 (0%)
kube-system          kube-controller-manager-ubuntu 200m (10%) 0 (0%)    0 (0%)    0 (0%)
kube-system          kube-dns-86f4d74b45-dz9gl 260m (13%) 0 (0%)    110Mi (5%) 170Mi (9%)
kube-system          kube-proxy-n9x4z      0 (0%)    0 (0%)    0 (0%)    0 (0%)
kube-system          kube-scheduler-ubuntu 100m (5%)  0 (0%)    0 (0%)    0 (0%)
kube-system          weave-net-kq9db       20m (1%)   0 (0%)    0 (0%)    0 (0%)
Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
CPU Requests  CPU Limits  Memory Requests  Memory Limits
-----
830m (41%)    0 (0%)      110Mi (5%)      170Mi (9%)
Events:
Type      Reason      Age   From      Message
----
Normal    NodeReady   37m   kubelet, ubuntu Node ubuntu status is now: NodeReady
user@ubuntu:~$

```

Describe provides a wealth of node information. Your report will be similar but different than the one above.

- How much memory does your node have?
- How many CPUs?
- How many pods can your node run?
- What container runtime is the `kubelet` using?
- What version of `kubelet` is your node running?

Previously we used the `version` subcommand to discover the version of the `kubect1` client but now that our config is in place we can also see the version of the cluster API Server.

```

user@ubuntu:~$ kubectl version

Client Version: version.Info{Major:"1", Minor:"10", GitVersion:"v1.10.0",
GitCommit:"fc32d2f3698e36b93322a3465f63a14e9f0eae", GitTreeState:"clean", BuildDate:"2018-03-26T16:55:54Z",
GoVersion:"go1.9.3", Compiler:"gc", Platform:"linux/amd64"}
Server Version: version.Info{Major:"1", Minor:"10", GitVersion:"v1.10.0",
GitCommit:"fc32d2f3698e36b93322a3465f63a14e9f0eae", GitTreeState:"clean", BuildDate:"2018-03-26T16:44:10Z",
GoVersion:"go1.9.3", Compiler:"gc", Platform:"linux/amd64"}
user@ubuntu:~$

```

If you are familiar with Golang, notice the use of the `gc` tool chain (vs `gccgo`).

3. Creating Applications

With our cluster running and `kubect1` configured we can try to start a simple application on the cluster. The `kubect1` command provides a `get` subcommand which can be used to get information on any one of the key Kubernetes component types: deployments, pods, replica sets, and Services. While you can type `kubect1 get replicaset`, that would be fairly inhumane so `kubect1` allows you to use the abbreviation `rs` for replica sets.

If you want to save yourself even more typing. Here is tab completion without the mentioned fix.

```

user@ubuntu:~$ kubectl get

Desktop/          .kube/          Public/
...

user@ubuntu:~$ kubectl get

```

You can enable temporary kubectl bash completion with:

```

user@ubuntu:~$ source <(kubectl completion bash)

user@ubuntu:~$

```

And after.

```

user@ubuntu:~$ kubectl get

certificatesigningrequest deployment          networkpolicy          replicaset
statefulset
...

user@ubuntu:~$ kubectl get

```

That is much better!

In a new shell, list the currently running services, deployments, replica sets, and pods on your cluster:

```
user@ubuntu:~$ kubectl get service,deployments,rs,pods
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	1h

```
user@ubuntu:~$
```

The only service running in our cluster is the *kubernetes* service itself. We have no deployments, replica sets, or pods yet (in our namespace). Do the same for the resources under the kube-system namespace, more on namespaces later.

```
user@ubuntu:~$ kubectl get service,deployments,rs,pods --namespace=kube-system
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP	1h

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
kube-dns	1	1	1	1	1h

NAME	DESIRED	CURRENT	READY	AGE
kube-dns-86f4d74b45	1	1	1	1h

NAME	READY	STATUS	RESTARTS	AGE
etcd-ubuntu	1/1	Running	0	1h
kube-apiserver-ubuntu	1/1	Running	0	1h
kube-controller-manager-ubuntu	1/1	Running	0	1h
kube-dns-86f4d74b45-dz9gl	3/3	Running	0	1h
kube-proxy-n9x4z	1/1	Running	0	1h
kube-scheduler-ubuntu	1/1	Running	0	1h
weave-net-kq9db	2/2	Running	0	48m

```
user@ubuntu:~$
```

We can view all namespaces via `--all-namespaces` (if we have permission).

To test our cluster lets run a single container pod with a replication factor of 2. When configured with the Docker Engine as the container manager, we can run any container image that Docker has preinstalled or knows how to download.

```
user@ubuntu:~$ kubectl run my-nginx --image=nginx:1.11 --replicas=2 --port=80
```

```
deployment.apps "my-nginx" created
user@ubuntu:~$
```

This creates two copies (replicas) of a pod running nginx. The deployment name is "my-nginx" and the image we used is "nginx", an official image pulled from Docker Hub by the Docker Engine in the background. The port switch tells Kubernetes the service port for our pod which will allow us to share the service with its users over that port (the program must actually use that port for this to work).

List the deployments running on the cluster:

```
user@ubuntu:~$ kubectl get deployment
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
my-nginx	2	2	2	0	7s

```
user@ubuntu:~$
```

This shows that our pods are deployed and up to date. It may take a bit to pull the Docker images (Available might be 0). Let's look at the pod listing from the Kubernetes side:

```
user@ubuntu:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
my-nginx-87464966f-7c6xj	1/1	Running	0	9s
my-nginx-87464966f-kdqqz	1/1	Running	0	9s

```
user@ubuntu:~$
```

You can use the `docker container ls` subcommand to display the containers running under the Docker Engine:

```
user@ubuntu:~$ docker container ls --filter "name=nginx"
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
PORTS	NAMES			

```

d91b69f1cb49      5766334bdaa0      "nginx -g 'daemon of..." 32 seconds ago      Up 32 seconds
k8s_my-nginx_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_0
292e9faf782e      5766334bdaa0      "nginx -g 'daemon of..." 32 seconds ago      Up 32 seconds
k8s_my-nginx_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0
3cd2b943db6b      k8s.gcr.io/pause-amd64:3.1  "/pause"              33 seconds ago      Up 32 seconds
k8s_POD_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0
3362db7d71a2      k8s.gcr.io/pause-amd64:3.1  "/pause"              33 seconds ago      Up 32 seconds
k8s_POD_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_
user@ubuntu:~$

```

As you can see, while our `run` subcommand requested that Kubernetes run `nginx` with 2 replicas, 4 containers were launched at that time.

In Kubernetes, each Pod instance has an infrastructure container, which is the first container that the `kubelet` instantiates. The infrastructure container uses the image `"gcr.io/google_containers/pause-amd64:3.1"` and acquires the pod's IP as well as a pod wide network and IPC namespace. All of the other containers in the pod then join the infrastructure container's network (`--net`) and IPC (`--ipc`) namespace allowing containers in the pod to easily communicate. The initial process (`"/pause"`) that runs in the infrastructure container does nothing, its sole purpose is to act as the anchor for the pod and its shared namespaces.

You can learn more about the pause container by looking at the source and ultimately what is `"pause()`".

- <https://github.com/kubernetes/kubernetes/tree/master/build/pause>
- <https://github.com/kubernetes/kubernetes/blob/master/build/pause/pause.c>
- `man 2 pause` or <http://man7.org/linux/man-pages/man2/pause.2.html>

The Docker listing shows us 4 containers, each pod having an infrastructure container (pause) and the container we asked for (nginx).

Kubernetes gives each pod a name and reports on the pod status, the number of times the pod has been restarted and the pod's uptime. You can find the pod names embedded in the container names displayed by the `docker container ls` command:

```

user@ubuntu:~$ docker container ls --filter "name=nginx" --format "{{.Names}}"

k8s_my-nginx_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_0
k8s_my-nginx_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0
k8s_POD_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0
k8s_POD_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_0
user@ubuntu:~$

```

Next let's ask Kubernetes to display the replica sets at work:

```

user@ubuntu:~$ kubectl get rs

NAME                DESIRED   CURRENT   READY   AGE
my-nginx-87464966f  2         2         2       7m
user@ubuntu:~$

```

As you can see, the deployment generated a replica set to create and manage our two pods. The replica set has the name we provided on the run command line above (plus numerical extension) and shows the number of replicas of the pod it is maintaining.

Try killing one of the nginx containers using the `docker container kill` subcommand and the ID of the underlying container based on the nginx image.

```

user@ubuntu:~$ docker container kill \
$(docker container ls --filter "ancestor=nginx:1.11" --format {{.ID}} | head -1)

d91b69f1cb49
user@ubuntu:~$

```

```

user@ubuntu:~$ docker container ls --filter "name=nginx"

CONTAINER ID        IMAGE               COMMAND              CREATED             STATUS
PORTS              NAMES
244e9e3849af        5766334bdaa0       "nginx -g 'daemon of..." 22 seconds ago      Up 21 seconds
k8s_my-nginx_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_1
292e9faf782e        5766334bdaa0       "nginx -g 'daemon of..." About a minute ago   Up About a minute
k8s_my-nginx_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0
3cd2b943db6b        k8s.gcr.io/pause-amd64:3.1  "/pause"            2 minutes ago       Up About a minute
k8s_POD_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0
3362db7d71a2        k8s.gcr.io/pause-amd64:3.1  "/pause"            2 minutes ago       Up About a minute
k8s_POD_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_0
user@ubuntu:~$

```

We can tell by the created time we have a new container. If you were fast enough, you may have seen the previous container exited. Docker terminates the container specified but Kubernetes has no knowledge of this action. When the Kubelet process, responsible for the pods

assigned to this node, sees the missing container, it simply reruns the nginx image.

After some time, if you run the previous command with the `-a` flag, we can see the previous killed container and the newly created one.

```
user@ubuntu:~$ docker container ls -a --filter "name=nginx"
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
244e9e3849af	5766334bdaa0	"nginx -g 'daemon of..."	43 seconds ago	Up 42 seconds
k8s_my-nginx_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_1				
d91b69f1cb49	5766334bdaa0	"nginx -g 'daemon of..."	2 minutes ago	Exited (137) 43 seconds ago
000c292e8a04_0		k8s_my-nginx_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-		
292e9faf782e	5766334bdaa0	"nginx -g 'daemon of..."	2 minutes ago	Up 2 minutes
k8s_my-nginx_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0				
3cd2b943db6b	k8s.gcr.io/pause-amd64:3.1	"/pause"	2 minutes ago	Up 2 minutes
k8s_POD_my-nginx-87464966f-kdqqz_default_c00da53b-32d6-11e8-885a-000c292e8a04_0				
3362db7d71a2	k8s.gcr.io/pause-amd64:3.1	"/pause"	2 minutes ago	Up 2 minutes
k8s_POD_my-nginx-87464966f-7c6xj_default_c00d0c6e-32d6-11e8-885a-000c292e8a04_0				

```
user@ubuntu:~$
```

Notice that we killed container `d91b69f1cb49` in the example but the new container `244e9e3849af` was created to replace it. Kubernetes does not "resurrect" containers that have failed. This is important because the container's state may be the reason it failed. Rather, Kubernetes runs a fresh copy of the original image, ensuring the container has a clean new internal state (cattle not pets!).

4. Create a Service

In modern software engineering terms, a service is an encapsulated set of functionality made available to consumers through an API. The problem with our nginx application at present is that it is replicated and when containers die new ones are created. The fact that there are multiple containers and that containers come and go makes using the app difficult.

To simplify things Kubernetes makes it possible for us to expose our dynamically replicated set of pods as a Service. The `kubectl` expose command does this.

Expose the my-nginx pod replica set as a service:

```
user@ubuntu:~$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
my-nginx-87464966f	2	2	2	8m

```
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl expose $(kubectl get rs -o=name) --port=80
```

```
service "my-nginx-87464966f" exposed
user@ubuntu:~$
```

This causes Kubernetes to create a conceptual Service for our pods, exposing the set of pods as a single endpoint for users. Use the `get services` subcommand to display your service.

```
user@ubuntu:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	2h
my-nginx-7cb86fff46	ClusterIP	10.102.46.131	<none>	80/TCP	21s

```
user@ubuntu:~$
```

Kubernetes has given our service a virtual IP (VIP) address and it will now distribute client connections across the pods running my-nginx.

To test the Service try curling it:

```
user@ubuntu:~$ NX_CIP=$(kubectl get services -o=custom-columns=NAME:.spec.clusterIP,NAME:.metadata.name \
| grep nginx | awk '{print $1}')
```

```
user@ubuntu:~$ echo $NX_CIP
10.102.46.131
```

```
user@ubuntu:~$ curl -I $NX_CIP
```

```
HTTP/1.1 200 OK
Server: nginx/1.11.13
Date: Wed, 28 Mar 2018 22:32:41 GMT
Content-Type: text/html
Content-Length: 612
```

```
Last-Modified: Tue, 04 Apr 2017 15:01:57 GMT
Connection: keep-alive
ETag: "58e3b565-264"
Accept-Ranges: bytes
user@ubuntu:~$
```

Which pod did you connect to? Does it matter?

5. Pod exec

While Kubernetes delegates all of the direct container operations to the container manager (usually Docker) it does pass through some useful container features.

For example, imagine you need to discover the distro of one of your pods' containers. You can use the `kubectl exec` subcommand to run arbitrary commands within a pod.

Try listing the running pods and then executing the `cat /etc/os-release` command within one of your pods.

```
user@ubuntu:~$ kubectl get pods

NAME                                READY   STATUS    RESTARTS   AGE
my-nginx-87464966f-7c6xj           1/1     Running   1           9m
my-nginx-87464966f-kdqqz           1/1     Running   0           9m
user@ubuntu:~$
```

Here we select the first one listed.

```
user@ubuntu:~$ NX_POD_NAME=$(kubectl get pods --no-headers | awk '{print $1}' | head -1)

user@ubuntu:~$ echo $NX_POD_NAME

my-nginx-87464966f-7c6xj
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl exec $NX_POD_NAME cat /etc/os-release

PRETTY_NAME="Debian GNU/Linux 8 (jessie)"
NAME="Debian GNU/Linux"
VERSION_ID="8"
VERSION="8 (jessie)"
ID=debian
HOME_URL="http://www.debian.org/"
SUPPORT_URL="http://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
user@ubuntu:~$
```

Running `cat /etc/os-release` via `kubectl exec` produces the information we needed. The `exec` subcommand chooses the first container within the pod to execute the command.

If you would like to execute the command within a specific container you can use the `-c` switch. The `describe pod` command will give you a list of the containers within the pod. We can also retrieve JSON output and filter for it. Our current pod has only one container but we can still test the command.

Try it:

```
user@ubuntu:~$ kubectl get pod $NX_POD_NAME -o json | jq ".spec.containers[].name" -r

my-nginx
user@ubuntu:~$
```

Use the `-c` switch to display the `os-release` file in the my-nginx container in the pod:

```
user@ubuntu:~$ kubectl exec -c my-nginx $NX_POD_NAME cat /etc/os-release

PRETTY_NAME="Debian GNU/Linux 8 (jessie)"
NAME="Debian GNU/Linux"
VERSION_ID="8"
VERSION="8 (jessie)"
ID=debian
HOME_URL="http://www.debian.org/"
SUPPORT_URL="http://www.debian.org/support"
BUG_REPORT_URL="https://bugs.debian.org/"
user@ubuntu:~$
```

Do not forget about tab completion!

6. System Logs

Each of the services composing our Kubernetes cluster emits a log file. In the current configuration, the `kubelet` log is controlled by systemd. You can use the `journalctl` command to tail (`-n`) the output for the kubelet service unit (`-u`)

```
user@ubuntu:~$ journalctl -n 400 --no-pager -u kubelet.service | grep -v "no observation"

-- Logs begin at Mon 2017-12-04 02:29:14 PST, end at Mon 2017-12-04 13:27:39 PST. --
Dec 04 13:04:27 ubuntu kubelet[74364]: E1204 13:04:27.551559 74364 helpers.go:468] PercpuUsage had 0 cpus, but the
actual number is 2; ignoring extra CPUs
Dec 04 13:04:27 ubuntu kubelet[74364]: E1204 13:04:27.566946 74364 helpers.go:468] PercpuUsage had 0 cpus, but the
actual number is 2; ignoring extra CPUs
Dec 04 13:04:27 ubuntu kubelet[74364]: I1204 13:04:27.609826 74364 reconciler.go:212]
operationExecutor.VerifyControllerAttachedVolume started for volume "default-token-2kmhb" (UniqueName:
"kubernetes.io/secret/b6bc1fa0-d936-11e7-a277-000c29ae8ddc-default-token-2kmhb") pod "my-nginx-7cb86fff46-b47g4"
(UUID: "b6bc1fa0-d936-11e7-a277-000c29ae8ddc")
Dec 04 13:04:27 ubuntu kubelet[74364]: I1204 13:04:27.710464 74364 reconciler.go:212]
operationExecutor.VerifyControllerAttachedVolume started for volume "default-token-2kmhb" (UniqueName:
"kubernetes.io/secret/b6be06ff-d936-11e7-a277-000c29ae8ddc-default-token-2kmhb") pod "my-nginx-7cb86fff46-pl8dc"
(UUID: "b6be06ff-d936-11e7-a277-000c29ae8ddc")
Dec 04 13:04:28 ubuntu kubelet[74364]: W1204 13:04:28.195920 74364 pod_container_deletor.go:77] Container
"9b38a47a31f8755eeb5c75e6b6c75fb735070adc862a6e76d80621f723011cc4" not found in pod's containers
Dec 04 13:04:28 ubuntu kubelet[74364]: W1204 13:04:28.228414 74364 pod_container_deletor.go:77] Container
"cf955e40c16f5f60d83b5663bfe5292c19f20297df543efed51e051c1c18db7c" not found in pod's containers
Dec 04 13:04:47 ubuntu kubelet[74364]: W1204 13:04:47.165440 74364 conversion.go:110] Could not get instant cpu
stats: different number of cpus
Dec 04 13:04:47 ubuntu kubelet[74364]: W1204 13:04:47.166876 74364 conversion.go:110] Could not get instant cpu
stats: different number of cpus
Dec 04 13:08:21 ubuntu kubelet[74364]: I1204 13:08:21.173299 74364 kuberuntime_manager.go:500] Container {Name:my-
nginx Image:nginx:1.11 Command:[] Args:[] WorkingDir: Ports:[{Name: HostPort:0 ContainerPort:80 Protocol:TCP
HostIP:}] EnvFrom:[] Env:[] Resources:{Limits:map[] Requests:map[]} VolumeMounts:[{Name:default-token-2kmhb
ReadOnly:true MountPath:/var/run/secrets/kubernetes.io/serviceaccount SubPath: MountPropagation:<nil>}]
LivenessProbe:nil ReadinessProbe:nil Lifecycle:nil TerminationMessagePath:/dev/termination-log
TerminationMessagePolicy:File ImagePullPolicy:IfNotPresent SecurityContext:nil Stdin:false StdinOnce:false
TTY:false} is dead, but RestartPolicy says that we should restart it.
Dec 04 13:08:21 ubuntu kubelet[74364]: I1204 13:08:21.173483 74364 kuberuntime_manager.go:739] checking backoff
for container "my-nginx" in pod "my-nginx-7cb86fff46-pl8dc_default(b6be06ff-d936-11e7-a277-000c29ae8ddc)"
user@ubuntu:~$
```

The rest of our services are running as containers.

We can use the `kubectl logs` command to display log output from our pods. Remember that Kubernetes system services run within the `kube-system` namespace by convention.

List the pods in the `kube-system` namespace:

```
user@ubuntu:~$ kubectl get pods --namespace=kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
etcd-ubuntu	1/1	Running	0	2h
kube-apiserver-ubuntu	1/1	Running	0	2h
kube-controller-manager-ubuntu	1/1	Running	0	2h
kube-dns-545bc4bfd4-957v6	3/3	Running	0	2h
kube-proxy-l9jv6	1/1	Running	0	2h
kube-scheduler-ubuntu	1/1	Running	0	2h
kubernetes-dashboard-7486b894c6-4bp24	1/1	Running	1	1h
weave-net-q2gds	2/2	Running	0	1h

```
user@ubuntu:~$
```

Now display the last 10 lines from the API service:

```
user@ubuntu:~$ kubectl logs --namespace=kube-system --tail=10 kube-apiserver-ubuntu
```

```
I1204 19:12:25.744624 1 storage_rbac.go:257] created role.rbac.authorization.k8s.io/system::leader-locking-
kube-scheduler in kube-system
I1204 19:12:25.785495 1 storage_rbac.go:257] created
role.rbac.authorization.k8s.io/system:controller:bootstrap-signer in kube-public
I1204 19:12:25.825925 1 storage_rbac.go:287] created rolebinding.rbac.authorization.k8s.io/system::leader-
locking-kube-controller-manager in kube-system
I1204 19:12:25.864200 1 storage_rbac.go:287] created rolebinding.rbac.authorization.k8s.io/system::leader-
locking-kube-scheduler in kube-system
I1204 19:12:25.904522 1 storage_rbac.go:287] created
rolebinding.rbac.authorization.k8s.io/system:controller:bootstrap-signer in kube-system
```

```
I1204 19:12:25.945576      1 storage_rbac.go:287] created
rolebinding.rbac.authorization.k8s.io/system:controller:cloud-provider in kube-system
I1204 19:12:25.986616      1 storage_rbac.go:287] created
rolebinding.rbac.authorization.k8s.io/system:controller:token-cleaner in kube-system
I1204 19:12:26.026056      1 storage_rbac.go:287] created
rolebinding.rbac.authorization.k8s.io/system:controller:bootstrap-signer in kube-public
E1204 21:25:32.974602      1 upgradeaware.go:310] Error proxying data from client to backend: write tcp
172.16.151.229:52660->172.16.151.229:10250: write: broken pipe
E1204 21:26:16.015490      1 upgradeaware.go:310] Error proxying data from client to backend: write tcp
172.16.151.229:52770->172.16.151.229:10250: write: broken pipe
user@ubuntu:~$
```

The log displays the timestamp of each entry, the API call being made and the service making the call.

- Examine the API log to see how often the kubelet reports in.

You can learn much by tracking the operations involved in starting a deployment.

Create a new single pod deployment with a descriptive name and then grep its activity in the logs.

```
user@ubuntu:~$ kubectl run mylogtracker --image nginx:1.11

deployment.apps "mylogtracker" created
user@ubuntu:~$
```

Again list the k8s system services, from here we can pick which logs to search for our new deployment.

```
user@ubuntu:~$ kubectl get pod --namespace=kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
etcd-ubuntu	1/1	Running	0	2h
kube-apiserver-ubuntu	1/1	Running	0	2h
kube-controller-manager-ubuntu	1/1	Running	0	2h
kube-dns-545bc4bfd4-957v6	3/3	Running	0	2h
kube-proxy-l9jv6	1/1	Running	0	2h
kube-scheduler-ubuntu	1/1	Running	0	2h
kubernetes-dashboard-7486b894c6-4bp24	1/1	Running	1	1h
weave-net-q2gds	2/2	Running	0	1h

```
user@ubuntu:~$
```

Try the controller manager server first:

```
user@ubuntu:~$ kubectl logs --namespace=kube-system kube-controller-manager-ubuntu | grep mylogtracker

I1204 21:28:34.996978      1 event.go:218] Event(v1.ObjectReference{Kind:"Deployment", Namespace:"default",
Name:"mylogtracker", UID:"157993cf-d93a-11e7-a277-000c29ae8ddc", APIVersion:"extensions", ResourceVersion:"9571",
FieldPath:""}): type: 'Normal' reason: 'ScalingReplicaSet' Scaled up replica set mylogtracker-6ff4ff6fd5 to 1
I1204 21:28:35.046040      1 event.go:218] Event(v1.ObjectReference{Kind:"ReplicaSet", Namespace:"default",
Name:"mylogtracker-6ff4ff6fd5", UID:"157bb1bc-d93a-11e7-a277-000c29ae8ddc", APIVersion:"extensions",
ResourceVersion:"9572", FieldPath:""}): type: 'Normal' reason: 'SuccessfulCreate' Created pod: mylogtracker-
6ff4ff6fd5-k5qrq
user@ubuntu:~$
```

You can track the interaction with the deployment, replica set, and pods through the API calls.

- How long does it take for the pod to become fully operational?

Now take a look at the kubelet log:

```
user@ubuntu:~$ journalctl -u kubelet.service | grep mylogtracker

Dec 04 13:28:35 ubuntu kubelet[74364]: I1204 13:28:35.182932 74364 reconciler.go:212]
operationExecutor.VerifyControllerAttachedVolume started for volume "default-token-2kmhb" (UniqueName:
"kubernetes.io/secret/1584a8f2-d93a-11e7-a277-000c29ae8ddc-default-token-2kmhb") pod "mylogtracker-6ff4ff6fd5-k5qrq"
(UID: "1584a8f2-d93a-11e7-a277-000c29ae8ddc")
user@ubuntu:~$
```

The first occurrence of our mylogtracker key shows the kubelet discovering the new pod during the **kubelet** sync loop with the API server (which you previously discovered runs every 10 seconds by default).

You can also view the events taking place within the Kubernetes cluster itself using the events resource type.

Try getting events with **kubectl** :

```
user@ubuntu:~$ kubectl get events
```

LAST SEEN TYPE	FIRST SEEN REASON	COUNT	NAME SOURCE	MESSAGE	KIND	SUBJECT
14m	14m	1	my-nginx-87464966f-7c6xj.152035d367cce61b	Pod		
Normal	Scheduled		default-scheduler	Successfully assigned my-nginx-87464966f-7c6xj to ubuntu		
14m	14m	1	my-nginx-87464966f-7c6xj.152035d37669e397	Pod		
Normal	SuccessfulMountVolume		kubelet, ubuntu	MountVolume.SetUp succeeded for volume "default-token-gggvw"		
12m	14m	2	my-nginx-87464966f-7c6xj.152035d3a3239349	Pod		spec.containers{my-nginx}
Normal	Pulled		kubelet, ubuntu	Container image "nginx:1.11" already present on machine		
12m	14m	2	my-nginx-87464966f-7c6xj.152035d3a9a03fb4	Pod		spec.containers{my-nginx}
Normal	Created		kubelet, ubuntu	Created container		
12m	14m	2	my-nginx-87464966f-7c6xj.152035d3b39fa51b	Pod		spec.containers{my-nginx}
Normal	Started		kubelet, ubuntu	Started container		
14m	14m	1	my-nginx-87464966f-gkqvq.152035cd1356f265	Pod		spec.containers{my-nginx}
Normal	Killing		kubelet, ubuntu	Killing container with id docker://my-nginx:Need to kill Pod		
14m	14m	1	my-nginx-87464966f-kdqqz.152035d36b657a88	Pod		
Normal	Scheduled		default-scheduler	Successfully assigned my-nginx-87464966f-kdqqz to ubuntu		
14m	14m	1	my-nginx-87464966f-kdqqz.152035d37658cf0c	Pod		
Normal	SuccessfulMountVolume		kubelet, ubuntu	MountVolume.SetUp succeeded for volume "default-token-gggvw"		
14m	14m	1	my-nginx-87464966f-kdqqz.152035d39ea7ffff	Pod		spec.containers{my-nginx}
Normal	Pulled		kubelet, ubuntu	Container image "nginx:1.11" already present on machine		
14m	14m	1	my-nginx-87464966f-kdqqz.152035d3a1f0507e	Pod		spec.containers{my-nginx}
Normal	Created		kubelet, ubuntu	Created container		
14m	14m	1	my-nginx-87464966f-kdqqz.152035d3aee00023	Pod		spec.containers{my-nginx}
Normal	Started		kubelet, ubuntu	Started container		
14m	14m	1	my-nginx-87464966f-tbpbk7.152035cd113d702a	Pod		spec.containers{my-nginx}
Normal	Killing		kubelet, ubuntu	Killing container with id docker://my-nginx:Need to kill Pod		
14m	14m	1	my-nginx-87464966f.152035cd06b132bb	ReplicaSet		
Normal	SuccessfulDelete		replicaset-controller	Deleted pod: my-nginx-87464966f-tbpbk7		
14m	14m	1	my-nginx-87464966f.152035cd0855e0b7	ReplicaSet		
Normal	SuccessfulDelete		replicaset-controller	Deleted pod: my-nginx-87464966f-gkqvq		
14m	14m	1	my-nginx-87464966f.152035d366ccaa19	ReplicaSet		
Normal	SuccessfulCreate		replicaset-controller	Created pod: my-nginx-87464966f-7c6xj		
14m	14m	1	my-nginx-87464966f.152035d367d79e18	ReplicaSet		
Normal	SuccessfulCreate		replicaset-controller	Created pod: my-nginx-87464966f-kdqqz		
14m	14m	1	my-nginx.152035cd05b19916	Deployment		
Normal	ScalingReplicaSet		deployment-controller	Scaled down replica set my-nginx-87464966f to 0		
14m	14m	1	my-nginx.152035d364e263a4	Deployment		
Normal	ScalingReplicaSet		deployment-controller	Scaled up replica set my-nginx-87464966f to 2		
1m	1m	1	mylogtracker-55d4cf8dfb-l676w.15203685a3eb9eb0	Pod		
Normal	Scheduled		default-scheduler	Successfully assigned mylogtracker-55d4cf8dfb-l676w to ubuntu		
1m	1m	1	mylogtracker-55d4cf8dfb-l676w.15203685b130b0ae	Pod		
Normal	SuccessfulMountVolume		kubelet, ubuntu	MountVolume.SetUp succeeded for volume "default-token-gggvw"		
1m	1m	1	mylogtracker-55d4cf8dfb-l676w.15203685cf91077b	Pod		spec.containers{mylogtracker}
Normal	Pulled		kubelet, ubuntu	Container image "nginx:1.11" already present on machine		
1m	1m	1	mylogtracker-55d4cf8dfb-l676w.15203685d2a89a10	Pod		spec.containers{mylogtracker}
Normal	Created		kubelet, ubuntu	Created container		
1m	1m	1	mylogtracker-55d4cf8dfb-l676w.15203685ded39d14	Pod		spec.containers{mylogtracker}
Normal	Started		kubelet, ubuntu	Started container		
1m	1m	1	mylogtracker-55d4cf8dfb.15203685a24fbee5	ReplicaSet		
Normal	SuccessfulCreate		replicaset-controller	Created pod: mylogtracker-55d4cf8dfb-l676w		
1m	1m	1	mylogtracker.15203685a074f9a6	Deployment		
Normal	ScalingReplicaSet		deployment-controller	Scaled up replica set mylogtracker-55d4cf8dfb to 1		

While your events will be different you can see the value of the cluster event log. You can display event data associated with a given resource by supplying its name. You can also control the output format.

For example to make the data machine readable you could output it in JSON:

```
user@ubuntu:~$ kubectl get -o json events | tail
    },
    "type": "Normal"
  }
],
"kind": "List",
```

```
"metadata": {
  "resourceVersion": "",
  "selfLink": ""
}
}
user@ubuntu:~$
```

7. Cleaning Up

Now that we have given our new cluster a good test we can clean up by deleting the service and deployments we have created. The `kubectl delete` subcommand allows you to delete objects you have created in the cluster.

To begin, delete the *my-nginx* Service:

```
user@ubuntu:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	5h
my-nginx-87464966f	ClusterIP	10.102.46.131	<none>	80/TCP	8m

```
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl delete service my-nginx-87464966f
```

```
service "my-nginx-87464966f" deleted
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	2h

```
user@ubuntu:~$
```

Do not delete the kubernetes service.

Next we can delete the deployments (which in turn removes the associated replica set).

```
user@ubuntu:~$ kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
my-nginx	2	2	2	2	23m
mylogtracker	1	1	1	1	4m

```
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl delete deployment my-nginx
```

```
deployment.extensions "my-nginx" deleted
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl delete deployment mylogtracker
```

```
deployment.extensions "mylogtracker" deleted
user@ubuntu:~$
```

```
user@ubuntu:~$ kubectl get deployments
```

```
No resources found.
user@ubuntu:~$
```

You Kubernetes cluster should now be cleaned up and ready for the next lab:

```
user@ubuntu:~$ kubectl get services,deployments,replicasets,pods
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
svc/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	2h

```
user@ubuntu:~$
```

Be sure to leave the `svc/kubernetes` services.

When we delete a deployment, Kubernetes ensures that all pods controlled by that replica set are also deleted. If you have trouble with deleting resources, the order matters. Deployments watch replica sets, replica sets watch pods; start with deployments. Services fall outside of that

restart logic, and can be deleted at any point.

Congratulations, you have completed the Kubernetes exploration lab!

Copyright (c) 2013-2018 RX-M LLC, Cloud Native Consulting, all rights reserved