

# Cloud Foundry Application Runtime

## Lab 13 – Running Docker Containers on CFAR

The Cloud Foundry Garden container engine supports OCI container images and can pull images from any system supporting the docker registry V2 API specification. In this lab we'll try running a docker container on cloud foundry and then exploring the container.

### 1. Run a public image on CFAR

To deploy a Docker image from a Docker Hub repository, run `cf push APP-NAME --docker-image REPO/IMAGE:TAG`. Replace the placeholder values in the command as follows:

- APP-NAME: The name of the app being pushed
- REPO: The name of the repository where the image is stored
- IMAGE: The name of an image from Docker Hub
- TAG: (Optional) The tag or version for the image

Push a standard nginx container from docker hub to your cloud foundry space (give the app a unique name):

```
user@ubuntu:~$ cf push my-app --docker-image nginx
Pushing app rxm-app33 to org rx-m.com / space development as randy.abernethy@rx-m.com...
Getting app info...
Creating app with these attributes...
+ name: rxm-app33
+ docker image: nginx
  routes:
+ rxm-app33.cfapps.io

Creating app rxm-app33...
Mapping routes...

Staging app and tracing logs...
  Cell a9e19260-c172-41da-a126-36e77b5ed4a8 successfully created container for instance 7f2f3118-0ada-4490-a541-e24d887be803
  Staging...
  Staging process started ...
  Staging process finished
  Exit status 0
  Staging Complete
  Cell a9e19260-c172-41da-a126-36e77b5ed4a8 stopping instance 7f2f3118-0ada-4490-a541-e24d887be803
  Cell a9e19260-c172-41da-a126-36e77b5ed4a8 destroying container for instance 7f2f3118-0ada-4490-a541-e24d887be803
  Cell a9e19260-c172-41da-a126-36e77b5ed4a8 successfully destroyed container for instance 7f2f3118-0ada-4490-a541-e24d887be803

Waiting for app to start...

name: rxm-app33
requested state: started
instances: 1/1
usage: 1G x 1 instances
routes: rxm-app33.cfapps.io
last uploaded: Fri 29 Jun 10:34:26 PDT 2018
stack: cflinuxfs2
docker image: nginx
start command: nginx -g daemon off;

      state   since                cpu    memory    disk      details
#0   running  2018-06-29T17:34:36Z    0.0%   0 of 1G   0 of 1G
```

Many Docker registries control access to images by authenticating with a username and password. You can add user and password authentication information to `cf push` if required:

- The `CF_DOCKER_PASSWORD` environment variable defines the password
- The `--docker-username` command line argument defines the username

### 2. Exploring the container

Let's use the Cloud Foundry `ssh` command to explore the contents of our new running container:

```
user@ubuntu:~$ cf ssh rxm-app33
```

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

The ssh command uses ssh to connect to the container. This can be helpful when debugging, inspection and monitoring processes.

List the files inside the container:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# ls -l /
total 8
drwxr-xr-x  2 root  root   4096 Jun 29 18:03 bin
drwxr-xr-x  2 root  root     6 Feb 23 23:23 boot
drwxr-xr-x  4 root  root   179 Jun 29 18:03 dev
drwxr-xr-x  1 root  root    97 Jun 29 18:03 etc
drwxr-xr-x  2 root  root     6 Feb 23 23:23 home
drwxr-xr-x  1 root  root    45 Jun 29 18:03 lib
drwxr-xr-x  2 root  root    34 Jun 29 18:03 lib64
drwxr-xr-x  2 root  root     6 Jun 25 00:00 media
drwxr-xr-x  2 root  root     6 Jun 25 00:00 mnt
drwxr-xr-x  2 root  root     6 Jun 25 00:00 opt
dr-xr-xr-x 871 nobody nogroup  0 Jun 29 18:03 proc
drwx----- 2 root  root    37 Jun 29 18:03 root
drwxr-xr-x  1 root  root    23 Jun 29 18:03 run
drwxr-xr-x  2 root  root  4096 Jun 29 18:03/sbin
drwxr-xr-x  2 root  root     6 Jun 25 00:00/srv
dr-xr-xr-x 13 nobody nogroup  0 Jun 23 08:02/sys
drwxrwxrwt  1 root  root    42 Jun 29 18:03 tmp
drwxr-xr-x  1 root  root    66 Jun 29 18:03/usr
drwxr-xr-x  1 root  root    19 Jun 29 18:03/var
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

The file system does not have a kernel! This is because it doesn't need one, application containers always use the kernel of the underlying host system.

Try listing all of the running processes:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# ps -ef
bash: ps: command not found
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

This container rootfs does not include a ps binary so we can not run the ps command (!). Instead try listing the files : in the standard Unix proc directory:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# ls -l /proc

total 0
dr-xr-xr-x  9 root  root      0 Jun 29 18:10 1
dr-xr-xr-x  9 root  root      0 Jun 29 18:10 102
dr-xr-xr-x  9 root  root      0 Jun 29 18:10 111
dr-xr-xr-x  9 root  root      0 Jun 29 18:10 35
dr-xr-xr-x  9 nginx nginx     0 Jun 29 18:10 50
dr-xr-xr-x  9 root  root      0 Jun 29 18:03 7
dr-xr-xr-x  9 root  root      0 Jun 29 18:10 71
dr-xr-xr-x  9 root  root      0 Jun 29 18:03 8
dr-xr-xr-x  9 root  root      0 Jun 29 18:03 89
dr-xr-xr-x  2 nobody nogroup  0 Jun 29 18:10/acpi
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/buddyinfo
dr-xr-xr-x  4 nobody nogroup  0 Jun 29 18:10/bus
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/cgroups
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/cmdline
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/consoles
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/cpuinfo
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/crypto
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/devices
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/diskstats
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/dma
dr-xr-xr-x  2 nobody nogroup  0 Jun 29 18:10/driver
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/execdomains
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/fb
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/filesystems
dr-xr-xr-x 10 nobody nogroup  0 Jun 29 18:10/fs
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/interrupts
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/iomem
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/ioports
dr-xr-xr-x 50 nobody nogroup  0 Jun 29 18:10/irq
-r--r--r--  1 nobody nogroup  0 Jun 29 18:10/kallsyms
```

```

crw-rw-rw- 1 nobody nogroup 1, 3 Jun 23 08:01 kcore
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 key-users
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 keys
-r----- 1 nobody nogroup 0 Jun 29 18:10 kmsg
-r----- 1 nobody nogroup 0 Jun 29 18:10 kpagecgroup
-r----- 1 nobody nogroup 0 Jun 29 18:10 kpagecount
-r----- 1 nobody nogroup 0 Jun 29 18:10 kpageflags
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 loadavg
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 locks
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 mdstat
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 meminfo
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 misc
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 modules
lrwxrwxrwx 1 nobody nogroup 11 Jun 29 18:10 mounts -> self/mounts
-rw-r--r-- 1 nobody nogroup 0 Jun 29 18:10 mtrr
lrwxrwxrwx 1 nobody nogroup 8 Jun 29 18:10 net -> self/net
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 pagetypeinfo
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 partitions
crw-rw-rw- 1 nobody nogroup 1, 3 Jun 23 08:01 sched_debug
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 schedstat
drwxrwxrwt 2 root root 40 Jun 29 18:03 scsi
lrwxrwxrwx 1 nobody nogroup 0 Jun 29 18:03 self -> 111
-r----- 1 nobody nogroup 0 Jun 29 18:10 slabinfo
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 softirqs
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 stat
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 swaps
dr-xr-xr-x 1 nobody nogroup 0 Jun 29 18:03 sys
--w----- 1 nobody nogroup 0 Jun 29 18:10 sysrq-trigger
dr-xr-xr-x 2 nobody nogroup 0 Jun 29 18:10 sysvipc
lrwxrwxrwx 1 nobody nogroup 0 Jun 29 18:03 thread-self -> 111/task/111
crw-rw-rw- 1 nobody nogroup 1, 3 Jun 23 08:01 timer_list
crw-rw-rw- 1 nobody nogroup 1, 3 Jun 23 08:01 timer_stats
dr-xr-xr-x 4 nobody nogroup 0 Jun 29 18:10 tty
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 uptime
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 version
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 version_signature
-r----- 1 nobody nogroup 0 Jun 29 18:10 vmallocinfo
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 vmstat
dr-xr-xr-x 2 nobody nogroup 0 Jun 29 18:10 xen
-r--r--r-- 1 nobody nogroup 0 Jun 29 18:10 zoneinfo

```

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

The /proc filesystem is used on Unix systems to give privileged users "file system" style access to system information. Note that most of the files have a size of 0. This is because they are not really files, they are simply filesystem constructs allowing us to query attributed of the system. All of the files with just a number for a name are actually process IDs for the running processes. List them:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# ls -l /proc/ | grep ' [0-9]*$'
```

```

total 0
dr-xr-xr-x 9 root root 0 Jun 29 18:18 1
dr-xr-xr-x 9 root root 0 Jun 29 18:18 102
dr-xr-xr-x 9 root root 0 Jun 29 18:18 135
dr-xr-xr-x 9 root root 0 Jun 29 18:18 136
dr-xr-xr-x 9 root root 0 Jun 29 18:18 35
dr-xr-xr-x 9 nginx nginx 0 Jun 29 18:18 50
dr-xr-xr-x 9 root root 0 Jun 29 18:03 7
dr-xr-xr-x 9 root root 0 Jun 29 18:18 71
dr-xr-xr-x 9 root root 0 Jun 29 18:03 8
dr-xr-xr-x 9 root root 0 Jun 29 18:03 89
lrwxrwxrwx 1 nobody nogroup 0 Jun 29 18:03 self -> 135

```

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

This command shows all of the files with a number for a name. The self link always points to the currently running process (which would have been the ls command in this example).

In Unix systems, process 1 is the first process executed after system boot or, in the case of containers with a process namespace, the first process to run in the container. List the contents of the "1" directory:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# ls -l /proc/1
```

```

total 0
dr-xr-xr-x 2 root root 0 Jun 29 18:20 attr

```

```

-rw-r--r-- 1 root root 0 Jun 29 18:20 autogroup
-r----- 1 root root 0 Jun 29 18:20 auxv
-r--r--r-- 1 root root 0 Jun 29 18:20 cgroup
--w----- 1 root root 0 Jun 29 18:20 clear_refs
-r--r--r-- 1 root root 0 Jun 29 18:20 cmdline
-rw-r--r-- 1 root root 0 Jun 29 18:20 comm
-rw-r--r-- 1 root root 0 Jun 29 18:20 coredump_filter
-r--r--r-- 1 root root 0 Jun 29 18:20 cpuset
lrwxrwxrwx 1 root root 0 Jun 29 18:20 cwd -> /
-r----- 1 root root 0 Jun 29 18:20 environ
lrwxrwxrwx 1 root root 0 Jun 29 18:20 exe -> /tmp/garden-init
dr-x----- 2 root root 0 Jun 29 18:20 fd
dr-x----- 2 root root 0 Jun 29 18:20 fdinfo
-rw-r--r-- 1 root root 0 Jun 29 18:20 gid_map
-r----- 1 root root 0 Jun 29 18:20 io
-r--r--r-- 1 root root 0 Jun 29 18:20 limits
-rw-r--r-- 1 root root 0 Jun 29 18:20 loginuid
dr-x----- 2 root root 0 Jun 29 18:20 map_files
-r--r--r-- 1 root root 0 Jun 29 18:20 maps
-rw----- 1 root root 0 Jun 29 18:20 mem
-r--r--r-- 1 root root 0 Jun 29 18:20 mountinfo
-r--r--r-- 1 root root 0 Jun 29 18:20 mounts
-r----- 1 root root 0 Jun 29 18:20 mountstats
dr-xr-xr-x 6 root root 0 Jun 29 18:20 net
dr-x--x--x 2 root root 0 Jun 29 18:20 ns
-r--r--r-- 1 root root 0 Jun 29 18:20 numa_maps
-rw-r--r-- 1 root root 0 Jun 29 18:20 oom_adj
-r--r--r-- 1 root root 0 Jun 29 18:20 oom_score
-rw-r--r-- 1 root root 0 Jun 29 18:20 oom_score_adj
-r----- 1 root root 0 Jun 29 18:20 pagemap
-r----- 1 root root 0 Jun 29 18:20 personality
-rw-r--r-- 1 root root 0 Jun 29 18:20 projid_map
lrwxrwxrwx 1 root root 0 Jun 29 18:20 root -> /
-rw-r--r-- 1 root root 0 Jun 29 18:20 sched
-r--r--r-- 1 root root 0 Jun 29 18:20 schedstat
-r--r--r-- 1 root root 0 Jun 29 18:20 sessionid
-rw-r--r-- 1 root root 0 Jun 29 18:20 setgroups
-r--r--r-- 1 root root 0 Jun 29 18:20 smaps
-r----- 1 root root 0 Jun 29 18:20 stack
-r--r--r-- 1 root root 0 Jun 29 18:20 stat
-r--r--r-- 1 root root 0 Jun 29 18:20 statm
-r--r--r-- 1 root root 0 Jun 29 18:20 status
-r----- 1 root root 0 Jun 29 18:20 syscall
dr-xr-xr-x 3 root root 0 Jun 29 18:20 task
-r--r--r-- 1 root root 0 Jun 29 18:20 timers
-rw-r--r-- 1 root root 0 Jun 29 18:20 uid_map
-r--r--r-- 1 root root 0 Jun 29 18:20 wchan

```

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

Each of these files describes some attribute of the PID 1 process. To display the command line the process is running cat the contents of the cmdline file:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# cat /proc/1/cmdline
```

```
/tmp/garden-init
```

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

The CF Garden container manager launches all containers with an initialization process called garden-init. Unix systems create processes with sequential process ids. Some processes exit right away leaving gaps in the id sequence. List the command line of the next sequential process id after 1:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# cat /proc/7/cmdline
```

```
/tmp/lifecycle/diego-sshd--allowedKeyExchanges=--address=0.0.0.0:2222--allowUnauthenticatedClients=false--
inheritDaemonEnv=true--allowedCiphers=--allowedMACs=--LogLevel=fatal--debugAddr=
```

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

This is the Diego sshd daemon that we are connected to. List the next process:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# cat /proc/8/cmdline
```

```
nginx: master process /usr/sbin/nginx -g daemon off;
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~#
```

This is the actual container image process we intended to run.  
Explore some more if you like, then exit the ssh shell:

```
root@2f3f1dfc-e9ec-4f2a-4d1e-085e:~# exit
exit
user@ubuntu:~$
```

## CHALLENGE (optional) Run a private image

In this challenge step you will push your trash can inventory image from MSA lab 3 to Docker Hub and then use `sf push` to deploy it to cloud foundry.

- Navigate to <https://hub.docker.com> using any browser.
- Create a free account (or login to your existing account)
- Push your trash can inventory image from MSA lab 3 to your private dockerhub account
  - Remember you can create names for images using "docker image tag", e.g. for an account named fredsmith and an image `inv` you would need to tag the image with a name like "fredsmith/inv" to push it to the fredsmith account on docker hub
- Now use "cf push" as per above to run your container in your PWS space
  - Remember you will need to add your name and password to the cf push command

Congratulations, you have completed the Lab.

*Copyright (c) 2013-2018 RX-M LLC, Cloud Native Consulting, all rights reserved*